

Pstat 131 HW1

Kalvin Goode (9454554), Patrick Chen (970890)

2019/4/11

```
algae<-read_table2("algaeBloom.txt",col_names=c('season','size','speed','mxPH','mn02',
'Cl','NO3','NH4','oPO4','PO4','Chla','a1','a2','a3','a4','a5','a6','a7'),na="XXXXXX")
glimpse(algae)
```

1.

(a)

```
algae %>%
  group_by(season) %>%
  dplyr::summarize(summary=n())
```

```
## # A tibble: 4 x 2
##   season summary
##   <chr>      <int>
## 1 autumn      40
## 2 spring      53
## 3 summer      45
## 4 winter      62
```

(b)

```
!all(!is.na(algae)) #If no missing value, is.na() shows all false, ! makes them all TRUE.
```

```
## [1] TRUE
```

```
#So, all() shows TRUE only if no missing value in data
```

Yes, there are missing values.

```
algae %>%
  dplyr::summarize_each(funs(mean(.,na.rm=TRUE)),c('mxPH':'Chla'))
```

```
## # A tibble: 1 x 8
##   mxPH mn02 Cl NO3 NH4 oPO4 PO4 Chla
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  8.01  9.12 43.6  3.28 501.  73.6 138.  14.0
```

```
algae %>%
  dplyr::summarize_each(funs(var(.,na.rm=TRUE)),c('mxPH':'Chla'))
```

```
## # A tibble: 1 x 8
##   mxPH mn02 Cl NO3 NH4 oPO4 PO4 Chla
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.358  5.72 2193. 14.3 3851585. 8306. 16639. 420.
```

We notice that among these chemicals, Cl, NH4, oPO4 and PO4 have large mean and variance ratio. Especially for NH4 and PO4, which have variances 100 times greater than their mean.

(c)

```
algae %>%
  dplyr::summarize_each(funs(median(.,na.rm=TRUE)),c('mxPH':'Chla'))
```

```
## # A tibble: 1 x 8
##   mxPH mn02    Cl   NO3   NH4  oP04   P04  Chla
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  8.06   9.8  32.7  2.68  103.  40.2  103.   5.48
```

```
algae %>%
  dplyr::summarize_each(funs(mad(., constant=1, na.rm=TRUE)), c('mxPH': 'Chla'))
```


```
## # A tibble: 1 x 8
##   mxPH mn02    Cl   NO3   NH4  oP04   P04  Chla
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 0.340  1.38  22.4  1.46  75.3  29.7  82.5   4.5
```

We see that the chemicals mxPH and mn02 have the largest difference in term of percentage between MAD and median. Other than this, we also see that MAD and median of that chemical is always smaller than mean and variance of that chemical. We believe that the reason is that MAD and median mitigate the effect on outliers for each chemical.

2.

(a)


```
ggplot(data=algae, aes(x=algae$mxPH, stat(density))) +
  geom_histogram() + ggtitle("mxPH")
```



From the graph above, we see that the graph is roughly normal and is not skewed.

(b)


```
ggplot(data=algae,aes(x=algae$mxPH))+geom_histogram(aes(y=stat(density)))+
  ggtitle("mxPH")+geom_density()+geom_rug()
```



(c)

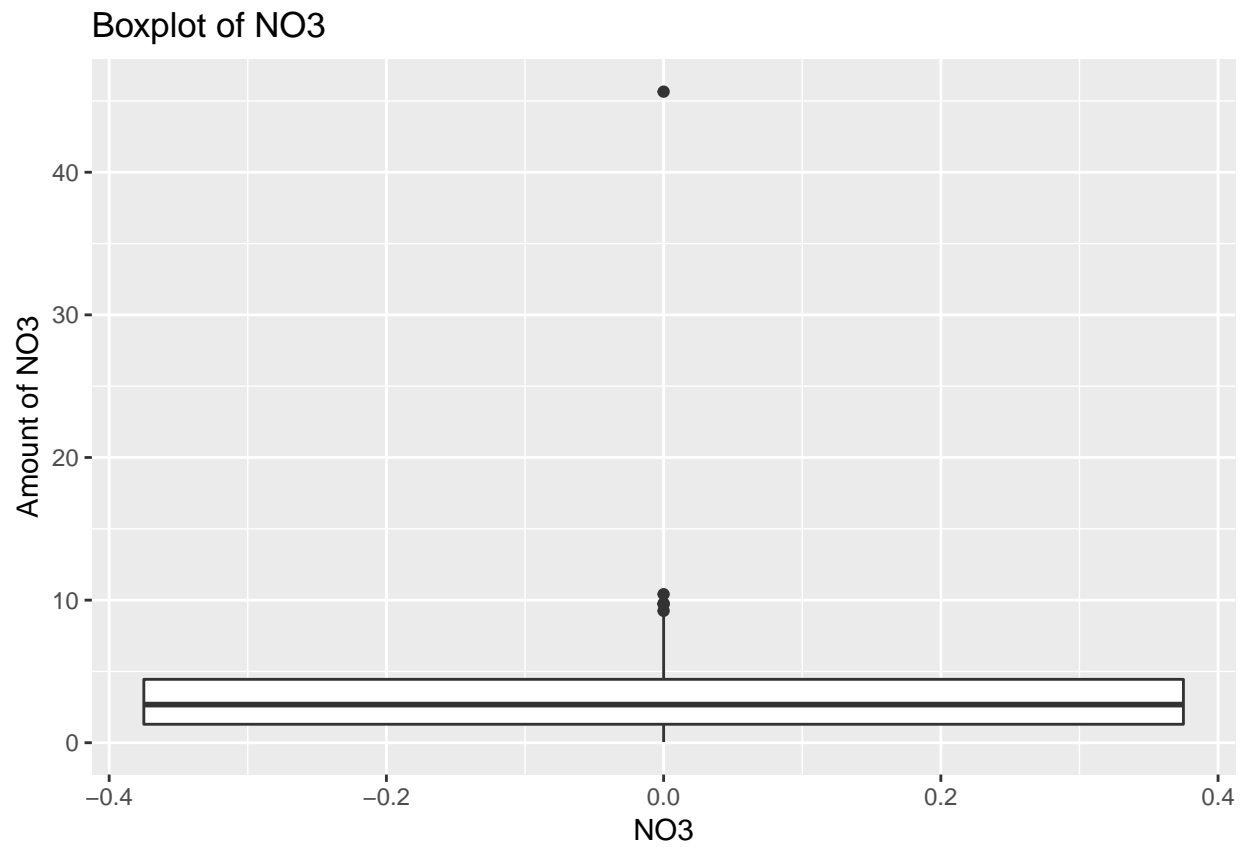
```
ggplot(data=algae,aes(x=algae$size,y=a1))+geom_boxplot()+
  xlab("a1")+ylab("Amount of a1")+
  ggtitle("A conditioned Boxplot of Algal a1")
```

A conditioned Boxplot of Algal a1

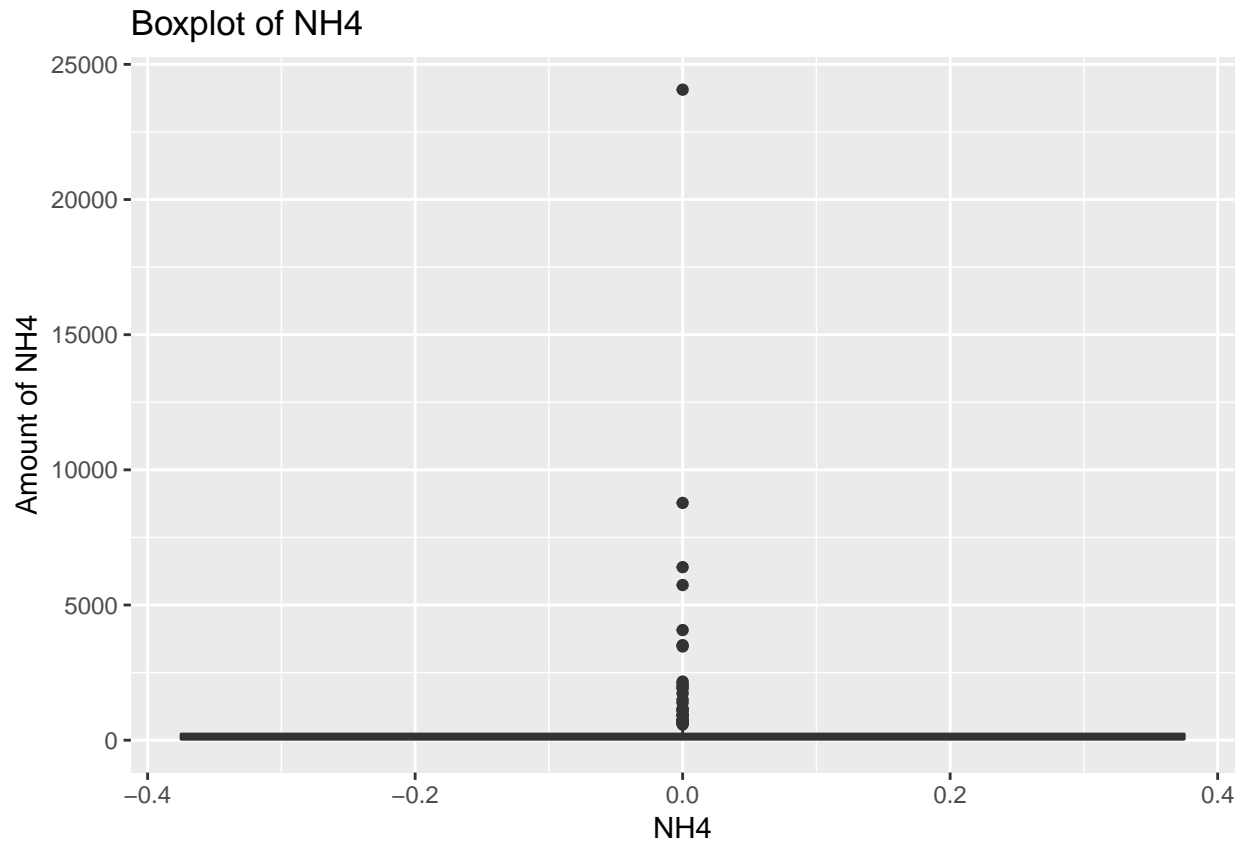


(d)

```
ggplot(data=algae,aes(,y=algae$N03))+geom_boxplot()+  
  xlab("N03")+ylab("Amount of N03")+  
  ggtitle("Boxplot of N03")
```



```
ggplot(data=algae,aes(,y=algae$NH4))+geom_boxplot()+  
  xlab("NH4")+ylab("Amount of NH4")+  
  ggtitle("Boxplot of NH4")
```



```

out=function(x)
{
  lower=quantile(x,0.25,na.rm=TRUE)-1.5*IQR(x,na.rm=TRUE)
  upper=quantile(x,0.75,na.rm=TRUE)+1.5*IQR(x,na.rm=TRUE)
  outliers=(x<lower)|(x>upper)
  count=length(na.omit(x[outliers]))
  return(count)
}

sprintf("The amount of outliers of NO3 and NH4 are %i and %i, respectively.",
        out(algae$NO3), out(algae$NH4))

```

```
## [1] "The amount of outliers of NO3 and NH4 are 5 and 27, respectively."
```

We show boxplot of NO3 and NH4 to check if outliers exists, then we wrote a function to count the amount of outlier for each elements. The defintion of an outlier is a number that is more than 1.5 times the length of data from lower or upper quartiles.

(e)

```

algae %>%
  dplyr::summarize_each(funs(mean(.,na.rm=TRUE),
                              var(.,na.rm=TRUE)), c("NO3","NH4"))

```

```

## # A tibble: 1 x 4
##   NO3_mean NH4_mean NO3_var  NH4_var
##   <dbl>     <dbl>   <dbl>   <dbl>
## 1     3.28     501.    14.3 3851585.

```

```
algae %>%
  dplyr::summarize_each(funs(median(.,na.rm=TRUE),
                               mad(.,constant=1,na.rm=TRUE)), c("NO3","NH4"))
```

```
## # A tibble: 1 x 4
##   NO3_median NH4_median NO3_mad NH4_mad
##       <dbl>       <dbl>   <dbl>   <dbl>
## 1       2.68       103.    1.46    75.3
```

For these two chemicals, we see that the ratio of mean and variance is significantly larger than the ratio of median and MAD. Therefore, we conclude that using median and MAD for measures is much more robust than using mean and variance when outliers are present. The outliers would dramatically increase the value of variance while having little effect on median and MAD.

3.

(a)

```
sprintf("There are %i variables and %i observations that have missing values.", sum(colSums(is.na(algae))
```

```
## [1] "There are 8 variables and 16 observations that have missing values."
```

```
sprintf("Below is the summary chart for amount of missing value for each predictors.")
```

```
## [1] "Below is the summary chart for amount of missing value for each predictors."
```

```
colSums(is.na(algae))
```

```
## season    size  speed  mxPH  mnO2    C1    NO3    NH4    oP04    P04
##      0      0      0      1      2    10     2      2      2      2
##   Chla    a1    a2    a3    a4    a5    a6    a7
##    12     0     0     0     0     0     0     0
```

(b)

```
algae.del=algae%>%
  filter(complete.cases(algae))
sprintf("Now, there are %i observation in algae.del.", nrow(algae.del))
```

```
## [1] "Now, there are 184 observation in algae.del."
```

(c)

```
algae.med=algae%>%
  mutate_at(vars(c("mxPH":"Chla")),
            ~ifelse(is.na(.),median(.,na.rm=TRUE),.))
sprintf("Now, there are %i observation in algae.med.",nrow(algae.med))
```

```
## [1] "Now, there are 200 observation in algae.med."
```

```
for (x in c(48,62,199))
{
  print(algae.med[x,4:11])
}
```

```
## # A tibble: 1 x 8
##   mxPH mnO2    C1    NO3    NH4    oP04    P04    Chla
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  8.06 12.6    9  0.23    10     5     6    1.1
## # A tibble: 1 x 8
##   mxPH mnO2    C1    NO3    NH4    oP04    P04    Chla
```

```
##      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      6.4   9.8  32.7  2.68 103.   40.2    14   5.48
## # A tibble: 1 x 8
##      mxPH  mn02    C1    N03   NH4   oP04    P04   Chla
##      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      8    7.6  32.7  2.68 103.   40.2  103.   5.48
```

(d)

```
cor(algae[,4:11],use="pairwise.complete.obs")
```

```
##              mxPH          mn02          C1          N03          NH4
## mxPH  1.00000000 -0.16861234  0.13610778 -0.13098054 -0.09353577
## mn02 -0.16861234  1.00000000 -0.27833251  0.09944373 -0.08747825
## C1    0.13610778 -0.27833251  1.00000000  0.22504091  0.07191298
## N03   -0.13098054  0.09944373  0.22504091  1.00000000  0.72144352
## NH4   -0.09353577 -0.08747825  0.07191298  0.72144352  1.00000000
## oP04  0.15899936 -0.41616295  0.39105351  0.14458782  0.22723723
## P04   0.18990814 -0.48748615  0.45744903  0.16860096  0.20818040
## Chla  0.44596182 -0.15326480  0.14985650  0.13967921  0.08894652
##              oP04          P04          Chla
## mxPH  0.1589994  0.1899081  0.44596182
## mn02 -0.4161629 -0.4874862 -0.15326480
## C1    0.3910535  0.4574490  0.14985650
## N03   0.1445878  0.1686010  0.13967921
## NH4   0.2272372  0.2081804  0.08894652
## oP04  1.0000000  0.9143652  0.11562132
## P04   0.9143652  1.0000000  0.25362130
## Chla  0.1156213  0.2536213  1.00000000
```

```
relation=lm(algae$P04~algae$oP04) #make a linear regression on y=P04 and x=oP04
predict(relation)[29]
```

```
##      30
## 58.41378
```

```
algae$P04[28]=predict(relation)[28]
algae$P04[28]
```

```
## [1] 76.51663
```

(e)

If we impute missing values with mean or medians, we may have a biased result because it would have smaller variance and does not reflect the uncertainty on prediction of unknown missing values. Such process would cause survivorship bias. Also, for small sample size, the estimation can be dramatically different compare with and without imputation.

4.

(a)

```
nfold=5
set.seed(10)
folds=cut(1:nrow(algae.med), breaks=nfold, labels=FALSE) %>% sample()
folds
```

```
##      [1] 3 2 3 4 1 2 2 2 3 3 4 3 1 3 2 2 1 2 2 4 4 3 4 2 2 5 4 2 4 2 3 1 1 4 2
##     [36] 5 4 5 5 3 2 1 1 3 1 1 4 2 5 4 5 4 4 2 1 5 5 4 2 5 1 5 2 2 3 3 2 2 1 5
##     [71] 5 2 5 4 4 4 2 5 1 1 3 2 1 5 1 4 3 1 1 5 4 2 2 4 2 3 4 3 2 1 1 1 4 1 3
```



```
## [106] 4 3 4 5 5 5 4 5 2 5 3 5 2 2 4 1 3 3 5 4 3 5 3 5 2 5 3 5 5 1 5 3 5 3 1
## [141] 4 1 4 3 5 5 1 1 4 2 4 4 2 1 4 5 4 1 3 2 4 4 4 1 5 5 3 5 3 5 3 3 5 1 3
## [176] 4 4 1 1 3 2 3 1 3 1 3 2 3 2 3 1 1 4 2 3 5 5 1 1 2
```

(b)

```
do.chunk <- function(chunkid, chunkdef, dat)
{ # function argument
  train = (chunkdef != chunkid)
  Xtr = dat[train,1:11] # get training set
  Ytr = dat[train,12] # get true response values in trainig set
  Xvl = dat[!train,1:11] # get validation set
  Yvl = dat[!train,12] # get true response values in validation set
  lm.a1 <- lm(a1~., data = dat[train,1:12])
  predYtr = predict(lm.a1) # predict training values
  predYvl = predict(lm.a1,Xvl) # predict validation values
  data.frame(fold = chunkid,
             train.error = mean((predYtr - Ytr$a1)^2),
             # compute and store training error
             val.error = mean((predYvl - Yvl$a1)^2))
             # compute and store test error
}

ldply(1:nfold, do.chunk, folds, algae.med)
```

```
##   fold train.error val.error
## 1    1    313.8769  183.6954
## 2    2    238.0351  552.8065
## 3    3    306.6797  229.1438
## 4    4    284.6941  325.9966
## 5    5    263.2112  410.4312
```

5.

(a)

```
algae.Test <- read_table2('algaeTest.txt',
col_names=c('season','size','speed','mxPH','mnO2','Cl','NO3',
'NH4','oPO4','PO4','Chla','a1'),
na=c('XXXXXX'))

lm.a1=lm(a1~., data=algae.med[,1:12])
Xtr = algae.med[,1:11] # get training set
Ytr = algae.med[,12] # get true response values in trainig set
Xvl = algae.Test[,1:11] # get validation set
Yvl = algae.Test[,12] # get true response values in validation set
lm.a1 <- lm(a1~., data = algae.med[,1:12])
predYtr = predict(lm.a1) # predict training values
predYvl = predict(lm.a1,Xvl) # predict validation values
train.error = mean((predYtr - Ytr$a1)^2) # compute and store training error
val.error = mean((predYvl - Yvl$a1)^2) # compute and store test error

train.error

## [1] 286.2661

val.error
```

```
## [1] 250.1794
```

We see that, compare with problem 4, the training error is roughly the same, but this new test error is smaller than most of the previous, test errors. This shows that this new model is fit for new dataset.

6.


(a)

```
library(ISLR)
head(Wage)
```

```
##      year age      maritl      race      education
## 231655 2006  18 1. Never Married 1. White      1. < HS Grad
## 86582  2004  24 1. Never Married 1. White      4. College Grad
## 161300 2003  45      2. Married 1. White      3. Some College
## 155159 2003  43      2. Married 3. Asian      4. College Grad
## 11443  2005  50      4. Divorced 1. White      2. HS Grad
## 376662 2008  54      2. Married 1. White      4. College Grad
##      region      jobclass      health health_ins
## 231655 2. Middle Atlantic 1. Industrial      1. <=Good      2. No
## 86582  2. Middle Atlantic 2. Information 2. >=Very Good      2. No
## 161300 2. Middle Atlantic 1. Industrial      1. <=Good      1. Yes
## 155159 2. Middle Atlantic 2. Information 2. >=Very Good      1. Yes
## 11443  2. Middle Atlantic 2. Information      1. <=Good      1. Yes
## 376662 2. Middle Atlantic 2. Information 2. >=Very Good      1. Yes
##      logwage      wage
## 231655 4.318063 75.04315
## 86582  4.255273 70.47602
## 161300 4.875061 130.98218
## 155159 5.041393 154.68529
## 11443  4.318063 75.04315
## 376662 4.845098 127.11574
```

```
data(Wage)
ggplot(data=Wage, aes(x=age, y=wage))+
  geom_point()+geom_smooth()+
  xlab("Age")+ylab("Wage")+
  ggtitle("Plot of wage vs age")
```

Plot of wage vs age



We see that the relationship between Age and Wage is not really linear. Notably between age 30 and 65, we see that there are outliers that earn above 250 while most of the wages is below 200. Also, as age increase, we see that there are less and less data points. These match our expectation because people tend to retire after 65 years old.

(b).

```
nfold=5
set.seed(11)
folds=cut(1:nrow(Wage), breaks=nfold, labels=FALSE) %>% sample()
do.chunk1=function(chunkid,folddef,dat) #constant polynomial
{
  train=(folddef!=chunkid)
  Xtr=dat[train,2]
  Ytr=dat[train,11]
  Xvl=dat[!train,2]
  Yvl=dat[!train,11]

  poly1=lm(wage~1,data=dat[train,c(2,11)])
  pdYtr=predict(poly1)
  pdYvl=predict(poly1,newdata=dat[!train,c(2,11)])
  data.frame(fold=chunkid,
    train.error=mean((pdYtr-Ytr)^2),
    # compute and store training error
    val.error=mean((pdYvl-Yvl)^2))
  # compute and store test error
}
```

```

table=NULL
table=rbind(table,ldply(1:nfold, do.chunk1,folds, Wage))
#combine result

do.chunk2=function(chunkid,folddef,dat,deg) #polynomial
{
  train=(folddef!=chunkid)
  Xtr=dat[train,2]
  Ytr=dat[train,11]
  Xvl=dat[!train,2]
  Yvl=dat[!train,11]

  poly=lm(wage~poly(age,degree=deg,raw=FALSE),
           data=dat[train,c(2,11)])
  pdYtr=predict(poly)
  pdYvl=predict(poly,newdata=dat[!train,c(2,11)])
  data.frame(fold=chunkid,
             train.error=mean((pdYtr-Ytr)^2),
             # compute and store training error
             val.error=mean((pdYvl-Yvl)^2))
             # compute and store test error
}

for(deg in 1:10)
  table=rbind(table,ldply(1:nfold, do.chunk2,folds, Wage, deg))

sequence=c(0,0,0,0,0,1,1,1,1,1,2,2,2,2,2,3,3,3,3,4,4,4,4,4,5,5,5,5,5,6,6,6,6,6,7,7,7,7,7,8,8,8,8,8,9,9,9,9,9,10,10,10,10,10)
DataFrame=cbind(table,sequence)
colnames(DataFrame)=c("fold","train.error","val.error","degree")
mean=DataFrame %>%
  group_by(degree) %>%
  summarize_at(.funs=funs(mean),.var=vars(train.error, val.error))

mean

```

```

## # A tibble: 11 x 3
##   degree train.error val.error
##   <dbl>      <dbl>      <dbl>
## 1     0        1740.        1743.
## 2     1        1674.        1679.
## 3     2        1597.        1603.
## 4     3        1592.        1598.
## 5     4        1590.        1596.
## 6     5        1589.        1597.
## 7     6        1588.        1596.
## 8     7        1587.        1597.
## 9     8        1587.        1598.
## 10    9        1584.        1597.
## 11   10        1584.        1597.

```


c.

```

plot(train.error~degree, data=mean, xlab="p", ylab="Training Error",
main="Plot of Training Error and Test Error",col="blue", type="l")

```

```
lines(mean$degree,mean$val.error,col="Red")
legend(7,1700,c("Training error", "Test error"), col=c("Blue","Red"), text.col=c("Blue","Red"),lty=c(1))
```



From the graph above, we see that the errors decrease dramatically between $p=0$ and $p=2$. After $p=2$, the test errors remain consistent. When choosing model, we want to have the most simplest model while having low test errors. Therefore, we decide that the model for this relationship is $wage = \beta_0 + \beta_1 age + \beta_2 age^2 + \epsilon$